

```
In [0]: 1 import tensorflow.compat.v1 as tf
        2 tf.disable_v2_behavior()
```

```
In [0]: 1 # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist\_cnn.py
        2
        3 from __future__ import print_function
        4 import keras
        5 from keras.datasets import mnist
        6 from keras.models import Sequential
        7 from keras.layers import Dense, Dropout, Flatten
        8 from keras.layers import Conv2D, MaxPooling2D
        9 from keras.layers.normalization import BatchNormalization
       10 from keras import backend as K
       11 from keras import regularizers
```

```
In [0]: 1 %matplotlib inline
        2 import matplotlib.pyplot as plt
        3 import numpy as np
        4 import time
        5 # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
        6 # https://stackoverflow.com/a/14434334
        7 # this function is used to update the plots for each epoch and error
        8 def plt_dynamic(x, vy, ty, ax, colors=['b']):
        9     ax.plot(x, ty, 'r', label="Train Loss")
       10     ax.plot(x, vy, 'b', label="Validation Loss")
       11     plt.legend()
       12     plt.grid()
       13     fig.canvas.draw()
       14     plt.show()
```

```
In [0]: 1 batch_size = 128
2 num_classes = 10
3 epochs = 12
4
5 # input image dimensions
6 img_rows, img_cols = 28, 28
7
8 # the data, split between train and test sets
9 (x_train, y_train), (x_test, y_test) = mnist.load_data()
10
11 if K.image_data_format() == 'channels_first':
12     x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
13     x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
14     input_shape = (1, img_rows, img_cols)
15 else:
16     x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
17     x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
18     input_shape = (img_rows, img_cols, 1)
19
20 x_train = x_train.astype('float32')
21 x_test = x_test.astype('float32')
22 x_train /= 255
23 x_test /= 255
24 print('x_train shape:', x_train.shape)
25 print(x_train.shape[0], 'train samples')
26 print(x_test.shape[0], 'test samples')
27
28 # convert class vectors to binary class matrices
29 y_train = keras.utils.to_categorical(y_train, num_classes)
30 y_test = keras.utils.to_categorical(y_test, num_classes)
31
```

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

#Model -I with 2 - Convnet layers

```

In [0]: 1 %%time
2 model = Sequential()
3 model.add(Conv2D(1, kernel_size=(2, 2), activation='relu',
4                 input_shape=input_shape, padding='same',
5                 kernel_initializer='he_normal'))
6 model.add(MaxPooling2D(pool_size=(2, 2)))
7 model.add(Dropout(0.25))
8
9 model.add(Conv2D(7, (2, 2), activation='relu',
10                padding='same', kernel_initializer='he_normal'))
11 model.add(MaxPooling2D(pool_size=(2, 2)))
12 model.add(Dropout(0.4))
13 model.add(Flatten())
14 model.add(Dense(16, activation='relu'))
15 model.add(BatchNormalization())
16 model.add(Dropout(0.5))
17 model.add(Dense(num_classes, activation='softmax'))
18
19 model.compile(loss=keras.losses.categorical_crossentropy,
20               optimizer=keras.optimizers.Adadelta(),
21               metrics=['accuracy'])
22
23 history = model.fit(x_train, y_train,
24                     batch_size=batch_size,
25                     epochs=epochs,
26                     verbose=1,
27                     validation_data=(x_test, y_test))
28 score = model.evaluate(x_test, y_test, verbose=0)
29 print('Test loss:', score[0])
30 print('Test accuracy:', score[1])
31
32 #Error plot
33 fig,ax = plt.subplots(1,1)
34 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
35
36 # List of epoch numbers
37 x = list(range(1,epochs+1))
38 vy = history.history['val_loss']
39 ty = history.history['loss']
40 plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12


60000/60000 [=====] - 21s 344us/step - loss: 1.6334 - acc: 0.4473 - val_loss: 1.0381 - val_acc: 0.7450

Epoch 2/12

```

60000/60000 [=====] - 18s 293us/step - loss: 1.1997 - acc: 0.5956 - val_loss: 0.8789 - val_acc: 0.7756
Epoch 3/12
60000/60000 [=====] - 18s 292us/step - loss: 1.1173 - acc: 0.6253 - val_loss: 0.7978 - val_acc: 0.7936
Epoch 4/12
60000/60000 [=====] - 18s 297us/step - loss: 1.0810 - acc: 0.6407 - val_loss: 0.7912 - val_acc: 0.7843
Epoch 5/12
60000/60000 [=====] - 18s 296us/step - loss: 1.0553 - acc: 0.6488 - val_loss: 0.6906 - val_acc: 0.8256
Epoch 6/12
60000/60000 [=====] - 18s 292us/step - loss: 1.0407 - acc: 0.6530 - val_loss: 0.7503 - val_acc: 0.7938
Epoch 7/12
60000/60000 [=====] - 18s 295us/step - loss: 1.0214 - acc: 0.6609 - val_loss: 0.7066 - val_acc: 0.8050
Epoch 8/12
60000/60000 [=====] - 18s 296us/step - loss: 1.0139 - acc: 0.6597 - val_loss: 0.6900 - val_acc: 0.8105
Epoch 9/12
60000/60000 [=====] - 18s 292us/step - loss: 0.9996 - acc: 0.6670 - val_loss: 0.6678 - val_acc: 0.8183
Epoch 10/12
60000/60000 [=====] - 18s 301us/step - loss: 0.9960 - acc: 0.6666 - val_loss: 0.6624 - val_acc: 0.8220
Epoch 11/12
60000/60000 [=====] - 17s 291us/step - loss: 0.9840 - acc: 0.6702 - val_loss: 0.7649 - val_acc: 0.7717
Epoch 12/12
60000/60000 [=====] - 18s 292us/step - loss: 0.9739 - acc: 0.6716 - val_loss: 0.7207 - val_acc: 0.7942
Test loss: 0.7206710824012756
Test accuracy: 0.7942

```



CPU times: user 5min 22s, sys: 25.6 s, total: 5min 48s
Wall time: 3min 37s

```

In [0]: 1 %%time
2 model = Sequential()
3 model.add(Conv2D(1, kernel_size=(2, 2), activation='relu',
4                 input_shape=input_shape, padding='same',
5                 kernel_initializer='he_normal'))
6 model.add(MaxPooling2D(pool_size=(2, 2)))
7 model.add(Dropout(0.25))
8
9 model.add(Conv2D(7, (2, 2), activation='relu',
10                padding='same', kernel_initializer='he_normal'))
11 model.add(MaxPooling2D(pool_size=(2, 2)))
12 model.add(Dropout(0.4))
13 model.add(Flatten())
14 model.add(Dense(16, activation='relu'))
15 model.add(BatchNormalization())
16 model.add(Dropout(0.5))
17 model.add(Dense(num_classes, activation='softmax', kernel_regularizer=regularizers.l2(0.01)))
18 model.compile(loss=keras.losses.categorical_crossentropy,
19               optimizer=keras.optimizers.Adadelta(),
20               metrics=['accuracy'])
21
22 history = model.fit(x_train, y_train,
23                    batch_size=batch_size,
24                    epochs=epochs,
25                    verbose=1,
26                    validation_data=(x_test, y_test))
27 score = model.evaluate(x_test, y_test, verbose=0)
28 print('Test loss:', score[0])
29 print('Test accuracy:', score[1])
30
31 #Error plot
32 fig,ax = plt.subplots(1,1)
33 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
34
35 # List of epoch numbers
36 x = list(range(1,epochs+1))
37 vy = history.history['val_loss']
38 ty = history.history['loss']
39 plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 20s 336us/step - loss: 1.9122 - acc: 0.3738 - val_loss: 1.0775 - val_acc: 0.7722

Epoch 2/12


60000/60000 [=====] - 17s 282us/step - loss: 1.4165 - acc: 0.5462 - val_loss: 0.8782 - val_acc: 0.8191

Epoch 3/12

```

60000/60000 [=====] - 17s 287us/step - loss: 1.2831 - acc: 0.5912 - val_loss: 0.8111 - val_acc: 0.8334
Epoch 4/12
60000/60000 [=====] - 17s 287us/step - loss: 1.2130 - acc: 0.6127 - val_loss: 0.7616 - val_acc: 0.8515
Epoch 5/12
60000/60000 [=====] - 17s 289us/step - loss: 1.1369 - acc: 0.6423 - val_loss: 0.7158 - val_acc: 0.8611
Epoch 6/12
60000/60000 [=====] - 17s 286us/step - loss: 1.0889 - acc: 0.6563 - val_loss: 0.6870 - val_acc: 0.8680
Epoch 7/12
60000/60000 [=====] - 17s 288us/step - loss: 1.0478 - acc: 0.6698 - val_loss: 0.6538 - val_acc: 0.8683
Epoch 8/12
60000/60000 [=====] - 17s 287us/step - loss: 1.0173 - acc: 0.6759 - val_loss: 0.6394 - val_acc: 0.8686
Epoch 9/12
60000/60000 [=====] - 17s 289us/step - loss: 0.9909 - acc: 0.6871 - val_loss: 0.6185 - val_acc: 0.8701
Epoch 10/12
60000/60000 [=====] - 17s 291us/step - loss: 0.9771 - acc: 0.6884 - val_loss: 0.6038 - val_acc: 0.8691
Epoch 11/12
60000/60000 [=====] - 17s 291us/step - loss: 0.9578 - acc: 0.6944 - val_loss: 0.6198 - val_acc: 0.8559
Epoch 12/12
60000/60000 [=====] - 17s 291us/step - loss: 0.9457 - acc: 0.6979 - val_loss: 0.6393 - val_acc: 0.8435
Test loss: 0.6393014723777771
Test accuracy: 0.8435

```



CPU times: user 5min 20s, sys: 21.6 s, total: 5min 42s
Wall time: 3min 33s

#Model -II with 3 - Convnet layers

In [0]:

```
1 %%time
2 model = Sequential()
3 model.add(Conv2D(16, kernel_size=(3, 3), activation='relu',
4                 input_shape=input_shape, padding='same',
5                 kernel_initializer='he_normal'))
6 model.add(MaxPooling2D(pool_size=(2, 2)))
7 model.add(Dropout(0.25))
8 model.add(Conv2D(32, (3, 3), activation='relu',
9                 padding='same', kernel_initializer='he_normal'))
10 model.add(MaxPooling2D(pool_size=(2, 2)))
11 model.add(BatchNormalization())
12 model.add(Dropout(0.5))
13
14 model.add(Conv2D(64, (3, 3), activation='relu',
15                 padding='same', kernel_initializer='he_normal'))
16 model.add(MaxPooling2D(pool_size=(2, 2)))
17 model.add(Dropout(0.4))
18 model.add(Flatten())
19 model.add(Dense(32, activation='relu'))
20 model.add(BatchNormalization())
21 model.add(Dropout(0.5))
22 model.add(Dense(num_classes, activation='softmax'))
23
24 model.compile(loss=keras.losses.categorical_crossentropy,
25               optimizer=keras.optimizers.Adadelta(),
26               metrics=['accuracy'])
27
28 history = model.fit(x_train, y_train,
29                     batch_size=batch_size,
30                     epochs=epochs,
31                     verbose=1,
32                     validation_data=(x_test, y_test))
33 score = model.evaluate(x_test, y_test, verbose=0)
34 print('Test loss:', score[0])
35 print('Test accuracy:', score[1])
36
37 #Error plot
38 fig,ax = plt.subplots(1,1)
39 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
40
41 # List of epoch numbers
42 x = list(range(1,epochs+1))
43 vy = history.history['val_loss']
44 ty = history.history['loss']
45 plt_dynamic(x, vy, ty, ax)
```


Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 55s 923us/step - loss: 1.3213 - acc: 0.5683 - val_loss: 0.3741 - val_acc: 0.9137

Epoch 2/12

60000/60000 [=====] - 53s 882us/step - loss: 0.5147 - acc: 0.8425 - val_loss: 0.1402 - val_acc: 0.9607

Epoch 3/12

60000/60000 [=====] - 53s 881us/step - loss: 0.3393 - acc: 0.8978 - val_loss: 0.0842 - val_acc: 0.9752

Epoch 4/12

60000/60000 [=====] - 53s 879us/step - loss: 0.2666 - acc: 0.9209 - val_loss: 0.0654 - val_acc: 0.9800

Epoch 5/12

60000/60000 [=====] - 53s 880us/step - loss: 0.2294 - acc: 0.9323 - val_loss: 0.0546 - val_acc: 0.9833

Epoch 6/12

60000/60000 [=====] - 52s 871us/step - loss: 0.2034 - acc: 0.9407 - val_loss: 0.0501 - val_acc: 0.9839

Epoch 7/12

60000/60000 [=====] - 53s 878us/step - loss: 0.1856 - acc: 0.9458 - val_loss: 0.0453 - val_acc: 0.9844

Epoch 8/12

60000/60000 [=====] - 53s 885us/step - loss: 0.1740 - acc: 0.9494 - val_loss: 0.0393 - val_acc: 0.9866

Epoch 9/12

60000/60000 [=====] - 53s 886us/step - loss: 0.1655 - acc: 0.9525 - val_loss: 0.0384 - val_acc: 0.9880

Epoch 10/12

60000/60000 [=====] - 53s 879us/step - loss: 0.1602 - acc: 0.9543 - val_loss: 0.0385 - val_acc: 0.9871

Epoch 11/12


60000/60000 [=====] - 53s 883us/step - loss: 0.1544 - acc: 0.9561 - val_loss: 0.0377 - val_acc: 0.9875

Epoch 12/12

60000/60000 [=====] - 53s 884us/step - loss: 0.1479 - acc: 0.9570 - val_loss: 0.0326 - val_acc: 0.9895

Test loss: 0.0325662221849896

Test accuracy: 0.9895



CPU times: user 18min 28s, sys: 59.9 s, total: 19min 28s
Wall time: 10min 41s

In [0]:

```
1 %%time
2 model = Sequential()
3 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu',
4                 input_shape=input_shape, padding='same',
5                 kernel_initializer='he_normal'))
6 model.add(MaxPooling2D(pool_size=(2, 2)))
7 model.add(BatchNormalization())
8 model.add(Dropout(0.5))
9 model.add(Conv2D(64, (3, 3), activation='relu',
10                padding='same', kernel_initializer='he_normal'))
11 model.add(MaxPooling2D(pool_size=(2, 2)))
12 model.add(BatchNormalization())
13 model.add(Dropout(0.5))
14
15 model.add(Conv2D(64, (3, 3), activation='relu',
16                padding='same', kernel_initializer='he_normal'))
17 model.add(MaxPooling2D(pool_size=(2, 2)))
18 model.add(Dropout(0.5))
19 model.add(Flatten())
20 model.add(Dense(32, activation='relu'))
21 model.add(BatchNormalization())
22 model.add(Dropout(0.5))
23 model.add(Dense(num_classes, activation='softmax'))
24
25 model.compile(loss=keras.losses.categorical_crossentropy,
26               optimizer=keras.optimizers.Adadelta(),
27               metrics=['accuracy'])
28
29 history = model.fit(x_train, y_train,
30                     batch_size=batch_size,
31                     epochs=epochs,
32                     verbose=1,
33                     validation_data=(x_test, y_test))
34 score = model.evaluate(x_test, y_test, verbose=0)
35 print('Test loss:', score[0])
36 print('Test accuracy:', score[1])
37
38 #Error plot
39 fig,ax = plt.subplots(1,1)
40 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
41
42 # list of epoch numbers
43 x = list(range(1,epochs+1))
44 vy = history.history['val_loss']
45 ty = history.history['loss']
```

```
46 plt_dynamic(x, vy, ty, ax)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 172s 3ms/step - loss: 1.0956 - acc: 0.6502 - val_loss: 0.1769 - val_acc: 0.9527

Epoch 2/12

60000/60000 [=====] - 169s 3ms/step - loss: 0.3578 - acc: 0.8956 - val_loss: 0.0791 - val_acc: 0.9759

Epoch 3/12

60000/60000 [=====] - 169s 3ms/step - loss: 0.2398 - acc: 0.9299 - val_loss: 0.0552 - val_acc: 0.9822

Epoch 4/12

60000/60000 [=====] - 171s 3ms/step - loss: 0.1975 - acc: 0.9431 - val_loss: 0.0444 - val_acc: 0.9859

Epoch 5/12

60000/60000 [=====] - 170s 3ms/step - loss: 0.1687 - acc: 0.9511 - val_loss: 0.0388 - val_acc: 0.9864

Epoch 6/12

60000/60000 [=====] - 170s 3ms/step - loss: 0.1493 - acc: 0.9567 - val_loss: 0.0372 - val_acc: 0.9870

Epoch 7/12

60000/60000 [=====] - 172s 3ms/step - loss: 0.1361 - acc: 0.9612 - val_loss: 0.0339 - val_acc: 0.9887

Epoch 8/12

60000/60000 [=====] - 169s 3ms/step - loss: 0.1286 - acc: 0.9630 - val_loss: 0.0311 - val_acc: 0.9892

Epoch 9/12

60000/60000 [=====] - 171s 3ms/step - loss: 0.1242 - acc: 0.9647 - val_loss: 0.0297 - val_acc: 0.9895

Epoch 10/12

60000/60000 [=====] - 170s 3ms/step - loss: 0.1189 - acc: 0.9663 - val_loss: 0.0303 - val_acc: 0.9904

Epoch 11/12


60000/60000 [=====] - 172s 3ms/step - loss: 0.1146 - acc: 0.9676 - val_loss: 0.0281 - val_acc: 0.9910

Epoch 12/12

60000/60000 [=====] - 168s 3ms/step - loss: 0.1089 - acc: 0.9691 - val_loss: 0.0264 - val_acc: 0.9909

Test loss: 0.026352619710826548

Test accuracy: 0.9909



CPU times:user 1h 3min 26s, sys: 1min 33s, total: 1h 4min 59s

Wall time: 34min 12s

```
In [0]: 1 %%time
2 model = Sequential()
3 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu',
4                 input_shape=input_shape, padding='same',
5                 kernel_initializer='he_normal'))
6 model.add(MaxPooling2D(pool_size=(2, 2)))
7 model.add(BatchNormalization())
8 model.add(Dropout(0.5))
9 model.add(Conv2D(64, (3, 3), activation='relu',
10                padding='same', kernel_initializer='he_normal', kernel_regularizer=regularizers.l2(0.01)))
11 model.add(MaxPooling2D(pool_size=(2, 2)))
12 model.add(BatchNormalization())
13 model.add(Dropout(0.5))
14
15 model.add(Conv2D(64, (3, 3), activation='relu',
16                padding='same', kernel_initializer='he_normal', kernel_regularizer=regularizers.l2(0.01)))
17 model.add(MaxPooling2D(pool_size=(2, 2)))
18 model.add(Dropout(0.5))
19 model.add(Flatten())
20 model.add(Dense(32, activation='relu'))
21 model.add(BatchNormalization())
22 model.add(Dropout(0.5))
23 model.add(Dense(num_classes, activation='softmax', kernel_regularizer=regularizers.l2(0.01)))
24
25 model.compile(loss=keras.losses.categorical_crossentropy,
26               optimizer=keras.optimizers.Adadelta(),
27               metrics=['accuracy'])
28
29 history = model.fit(x_train, y_train,
30                    batch_size=batch_size,
31                    epochs=epochs,
32                    verbose=1,
33                    validation_data=(x_test, y_test))
34 score = model.evaluate(x_test, y_test, verbose=0)
35 print('Test loss:', score[0])
36 print('Test accuracy:', score[1])
37
38 #Error plot
39 fig,ax = plt.subplots(1,1)
40 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
41
42 # list of epoch numbers
43 x = list(range(1,epochs+1))
44 vy = history.history['val_loss']
45 ty = history.history['loss']
```

```
46 plt_dynamic(x, vy, ty, ax)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 172s 3ms/step - loss: 2.4369 - acc: 0.6701 - val_loss: 0.7775 - val_acc: 0.9349

Epoch 2/12

60000/60000 [=====] - 168s 3ms/step - loss: 0.5907 - acc: 0.9201 - val_loss: 0.3096 - val_acc: 0.9803

Epoch 3/12

60000/60000 [=====] - 170s 3ms/step - loss: 0.4379 - acc: 0.9408 - val_loss: 0.2606 - val_acc: 0.9849

Epoch 4/12

60000/60000 [=====] - 168s 3ms/step - loss: 0.3853 - acc: 0.9491 - val_loss: 0.2357 - val_acc: 0.9864

Epoch 5/12

60000/60000 [=====] - 168s 3ms/step - loss: 0.3562 - acc: 0.9537 - val_loss: 0.2255 - val_acc: 0.9835

Epoch 6/12

60000/60000 [=====] - 169s 3ms/step - loss: 0.3340 - acc: 0.9567 - val_loss: 0.2045 - val_acc: 0.9871

Epoch 7/12

60000/60000 [=====] - 168s 3ms/step - loss: 0.3210 - acc: 0.9572 - val_loss: 0.1992 - val_acc: 0.9868

Epoch 8/12

60000/60000 [=====] - 171s 3ms/step - loss: 0.3095 - acc: 0.9583 - val_loss: 0.1927 - val_acc: 0.9877

Epoch 9/12

60000/60000 [=====] - 168s 3ms/step - loss: 0.3017 - acc: 0.9602 - val_loss: 0.1872 - val_acc: 0.9871

Epoch 10/12

60000/60000 [=====] - 168s 3ms/step - loss: 0.2945 - acc: 0.9610 - val_loss: 0.1867 - val_acc: 0.9882

Epoch 11/12


60000/60000 [=====] - 167s 3ms/step - loss: 0.2871 - acc: 0.9618 - val_loss: 0.1834 - val_acc: 0.9881

Epoch 12/12

60000/60000 [=====] - 169s 3ms/step - loss: 0.2847 - acc: 0.9616 - val_loss: 0.1865 - val_acc: 0.9881

Test loss: 0.18646716620922088

Test accuracy: 0.9881



CPU times:user 1h 2min 51s, sys: 1min 34s, total: 1h 4min 25s

Wall time: 33min 55s

#Model - III with 5 - Convnet layers

In [0]:

```
1 %%time
2 model = Sequential()
3 model.add(Conv2D(2, kernel_size=(5, 5), activation='relu',
4                 input_shape=input_shape, padding='same',
5                 kernel_initializer='he_normal'))
6 model.add(MaxPooling2D(pool_size=(2, 2)))
7 model.add(Dropout(0.5))
8
9 model.add(Conv2D(4, (5, 5), activation='relu',
10                padding='same', kernel_initializer='he_normal'))
11 #model.add(MaxPooling2D(pool_size=(2, 2)))
12 model.add(BatchNormalization())
13
14 model.add(Conv2D(8, (5, 5), activation='relu',
15                padding='same', kernel_initializer='he_normal'))
16 model.add(MaxPooling2D(pool_size=(2, 2)))
17 model.add(BatchNormalization())
18 model.add(Dropout(0.5))
19
20 model.add(Conv2D(16, (5, 5), activation='relu',
21                padding='same', kernel_initializer='he_normal'))
22 model.add(MaxPooling2D(pool_size=(2, 2)))
23 model.add(BatchNormalization())
24 model.add(Dropout(0.5))
25
26 model.add(Conv2D(32, (5, 5), activation='relu',
27                padding='same', kernel_initializer='he_normal'))
28 model.add(MaxPooling2D(pool_size=(2, 2)))
29 model.add(Dropout(0.4))
30 model.add(Flatten())
31 model.add(Dense(32, activation='relu'))
32 model.add(BatchNormalization())
33 model.add(Dropout(0.5))
34 model.add(Dense(num_classes, activation='softmax'))
35
36 model.compile(loss=keras.losses.categorical_crossentropy,
37               optimizer=keras.optimizers.Adadelta(),
38               metrics=['accuracy'])
39
40 history = model.fit(x_train, y_train,
41                    batch_size=batch_size,
42                    epochs=epochs,
43                    verbose=1,
44                    validation_data=(x_test, y_test))
45 score = model.evaluate(x_test, y_test, verbose=0)
```

```

46 print('Test loss:', score[0])
47 print('Test accuracy:', score[1])
48
49 #Error plot
50 fig,ax = plt.subplots(1,1)
51 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
52
53 # List of epoch numbers
54 x = list(range(1,epochs+1))
55 vy = history.history['val_loss']
56 ty = history.history['loss']
57 plt_dynamic(x, vy, ty, ax)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 57s 947us/step - loss: 2.2582 - acc: 0.2415 - val_loss: 2.0836 - val_acc: 0.2941

Epoch 2/12

60000/60000 [=====] - 52s 873us/step - loss: 1.4995 - acc: 0.4536 - val_loss: 1.4567 - val_acc: 0.4895

Epoch 3/12

60000/60000 [=====] - 53s 885us/step - loss: 1.2549 - acc: 0.5537 - val_loss: 1.4632 - val_acc: 0.4759

Epoch 4/12

60000/60000 [=====] - 53s 886us/step - loss: 1.1002 - acc: 0.6265 - val_loss: 1.4953 - val_acc: 0.4831

Epoch 5/12

60000/60000 [=====] - 53s 884us/step - loss: 0.9788 - acc: 0.6817 - val_loss: 1.2909 - val_acc: 0.5163

Epoch 6/12

60000/60000 [=====] - 53s 881us/step - loss: 0.8790 - acc: 0.7172 - val_loss: 1.0550 - val_acc: 0.6089

Epoch 7/12

60000/60000 [=====] - 53s 887us/step - loss: 0.8037 - acc: 0.7490 - val_loss: 1.0653 - val_acc: 0.6219

Epoch 8/12

60000/60000 [=====] - 53s 891us/step - loss: 0.7418 - acc: 0.7761 - val_loss: 0.9224 - val_acc: 0.6829

Epoch 9/12

60000/60000 [=====] - 53s 891us/step - loss: 0.7057 - acc: 0.7873 - val_loss: 0.9815 - val_acc: 0.6719

Epoch 10/12

60000/60000 [=====] - 53s 890us/step - loss: 0.6689 - acc: 0.7990 - val_loss: 1.0064 - val_acc: 0.6757

Epoch 11/12


60000/60000 [=====] - 53s 887us/step - loss: 0.6529 - acc: 0.8068 - val_loss: 0.7009 - val_acc: 0.7508

Epoch 12/12

60000/60000 [=====] - 53s 887us/step - loss: 0.6297 - acc: 0.8162 - val_loss: 1.0190 - val_acc: 0.6963

Test loss: 1.0190036898136139

Test accuracy: 0.6963



CPU times: user 18min 52s, sys: 38.1 s, total: 19min 30s
Wall time: 10min 47s

Observation: This model is doing good with a relatively low loss and high accuracy and is not overfitting very much.

Summary

```
In [2]: 1 from prettytable import PrettyTable
2 pt = PrettyTable()
3 pt.field_names = ["S.No", "No.of Layers(Conv2D)", "Activation", "Kernel Size", "No of Channels", "Dropout", "BatchNormalization", "Max Pooling (2x2)", "L2-Regularizers", "Test loss", "Test accuracy"]
4 pt.add_row(["1", "3", "relu", "3X3", "(16,32,64)", "(0.25,0.5,0.4,0.5)", "Applied", "Applied", "Not Applied", "0.033", "0.989"])
5 pt.add_row(["2", "3", "relu", "3X3", "(64,64,64)", "(0.5,0.5,0.5,0.5)", "Applied", "Applied", "Not Applied", "0.026", "0.990"])
6 pt.add_row(["3", "3", "relu", "3X3", "(64,64,64)", "(0.5,0.5,0.5,0.5)", "Applied", "Applied", "Applied", "0.186", "0.981"])
7 pt.add_row(["4", "2", "relu", "2X2", "(1,7)", "(0.25,0.4,0.5)", "Applied", "Applied", "Not Applied", "0.720", "0.794"])
8 pt.add_row(["5", "2", "relu", "2X2", "(1,7)", "(0.25,0.4,0.5)", "Applied", "Applied", "Applied", "0.639", "0.843"])
9 pt.add_row(["6", "2", "relu", "2X2", "(1,7)", "(0.25,0.4,0.5)", "Applied", "Applied", "Applied", "1.116", "0.666"])
10 pt.add_row(["7", "5", "relu", "5X5", "(2,4,8,16,32)", "(0.5,0.5,0.5,0.4,0.5)", "Applied", "Applied", "Not Applied", "1.019", "0.696"])
11 print(pt)
```

S.No	No.of Layers(Conv2D)	Activation	Kernel Size	No of Channels	Dropout	BatchNormalization	Max Pooling (2x2)	L2-Regularizers	Test loss	Test accuracy
1	3	relu	3X3	(16,32,64)	(0.25,0.5,0.4,0.5)	Applied	Applied	Not Applied	0.033	0.989
2	3	relu	3X3	(64,64,64)	(0.5,0.5,0.5,0.5)	Applied	Applied	Not Applied	0.026	0.990
3	3	relu	3X3	(64,64,64)	(0.5,0.5,0.5,0.5)	Applied	Applied	Applied	0.186	0.981
4	2	relu	2X2	(1,7)	(0.25,0.4,0.5)	Applied	Applied	Not Applied	0.720	0.794
5	2	relu	2X2	(1,7)	(0.25,0.4,0.5)	Applied	Applied	Applied	0.639	0.843
6	2	relu	2X2	(1,7)	(0.25,0.4,0.5)	Applied	Applied	Applied	1.116	0.666
7	5	relu	5X5	(2,4,8,16,32)	(0.5,0.5,0.5,0.4,0.5)	Applied	Applied	Not Applied	1.019	0.696

Conclusion:

1. From the above table the model with 3 Convolution layers and 3x3x64 Kernel filters work very well with the Test loss of 0.026 and test accuracy 0.990.
2. After applying the L2 regularizer though it increased the test loss slightly and reduced the test accuracy a bit, it made the model to avoid overfit for some extent.
3. Techniques like Dropout, Maxpooling and Batch normalization plays a vital role in the model performance.
4. It is better to have the kernal dimension in odd numbers.

5. Since the image dimension is small, a kernel size of 5×5 is sufficient. But when the image dimension gets higher we can even use kernels with dimension upto 7×7 or more.